

Technical and Design Document: RAG-v1 Retrieval-Augmented Generation System

Abstract

RAG-v1 is a modular, production-ready Retrieval-Augmented Generation (RAG) system for document-based question answering. It integrates ChromaDB for vector search, LangChain for orchestration, and GROQ for large language model (LLM) inference. The system supports both a modern Streamlit web interface and a comprehensive command-line interface (CLI), enabling flexible, scalable, and secure document ingestion, retrieval, and conversational AI workflows.

1. Introduction and Motivation

The exponential growth of unstructured data has created a need for systems that can efficiently retrieve and synthesize information from large document collections. RAG-v1 addresses this by combining state-of-the-art vector search with LLMs, providing both technical and non-technical users with powerful tools for document-based Q&A, analytics, and knowledge management.

2. System Architecture

2.1 High-Level Overview

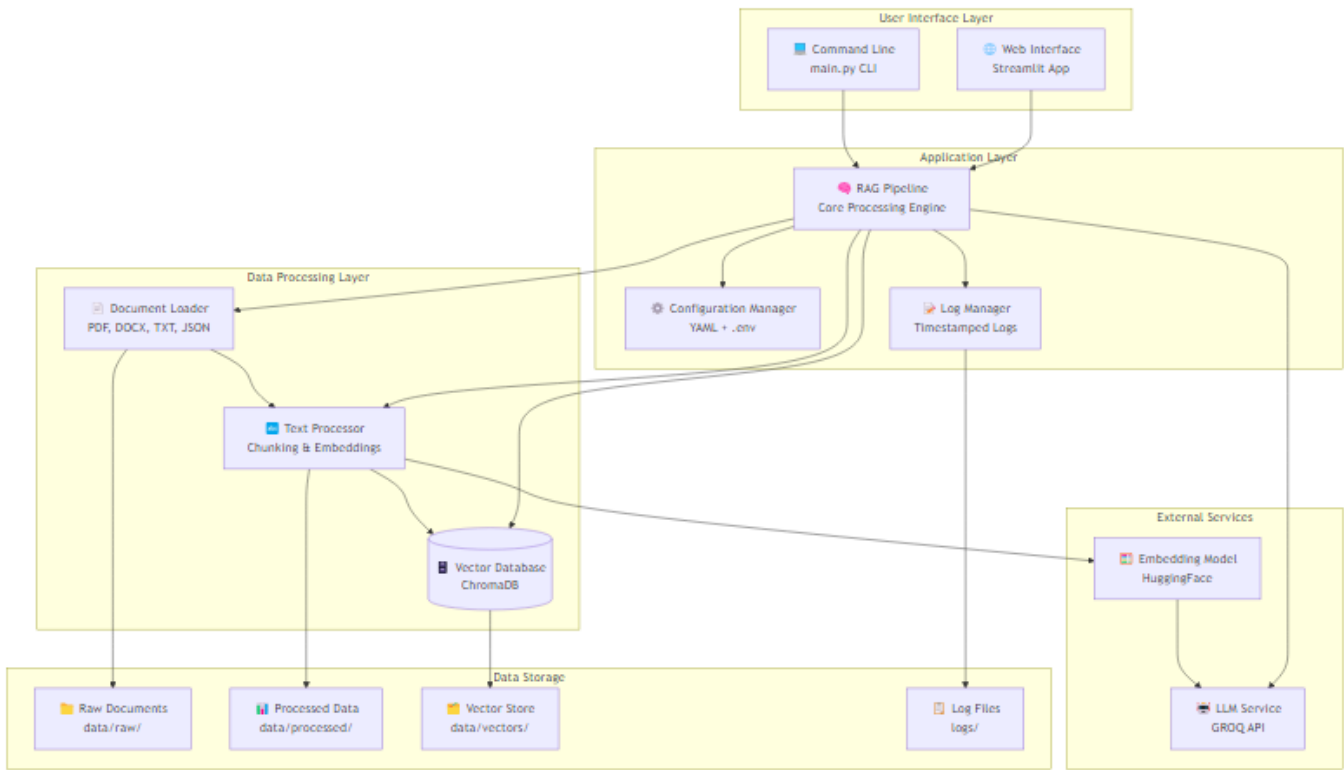


Figure 1: High-level system architecture showing all major components and their relationships

The RAG-v1 system is organized into five main layers:

- **User Interface Layer:** Dual interface—Streamlit web UI and CLI—for all operations.
- **Application Layer:** Core processing engine with configuration and logging management.

- **Data Processing Layer:** Document loading, text processing, and vector database operations.
- **External Services:** LLM and embedding model integrations.
- **Data Storage:** File system organization for raw documents, processed data, vectors, and logs.

2.2 Component Diagram

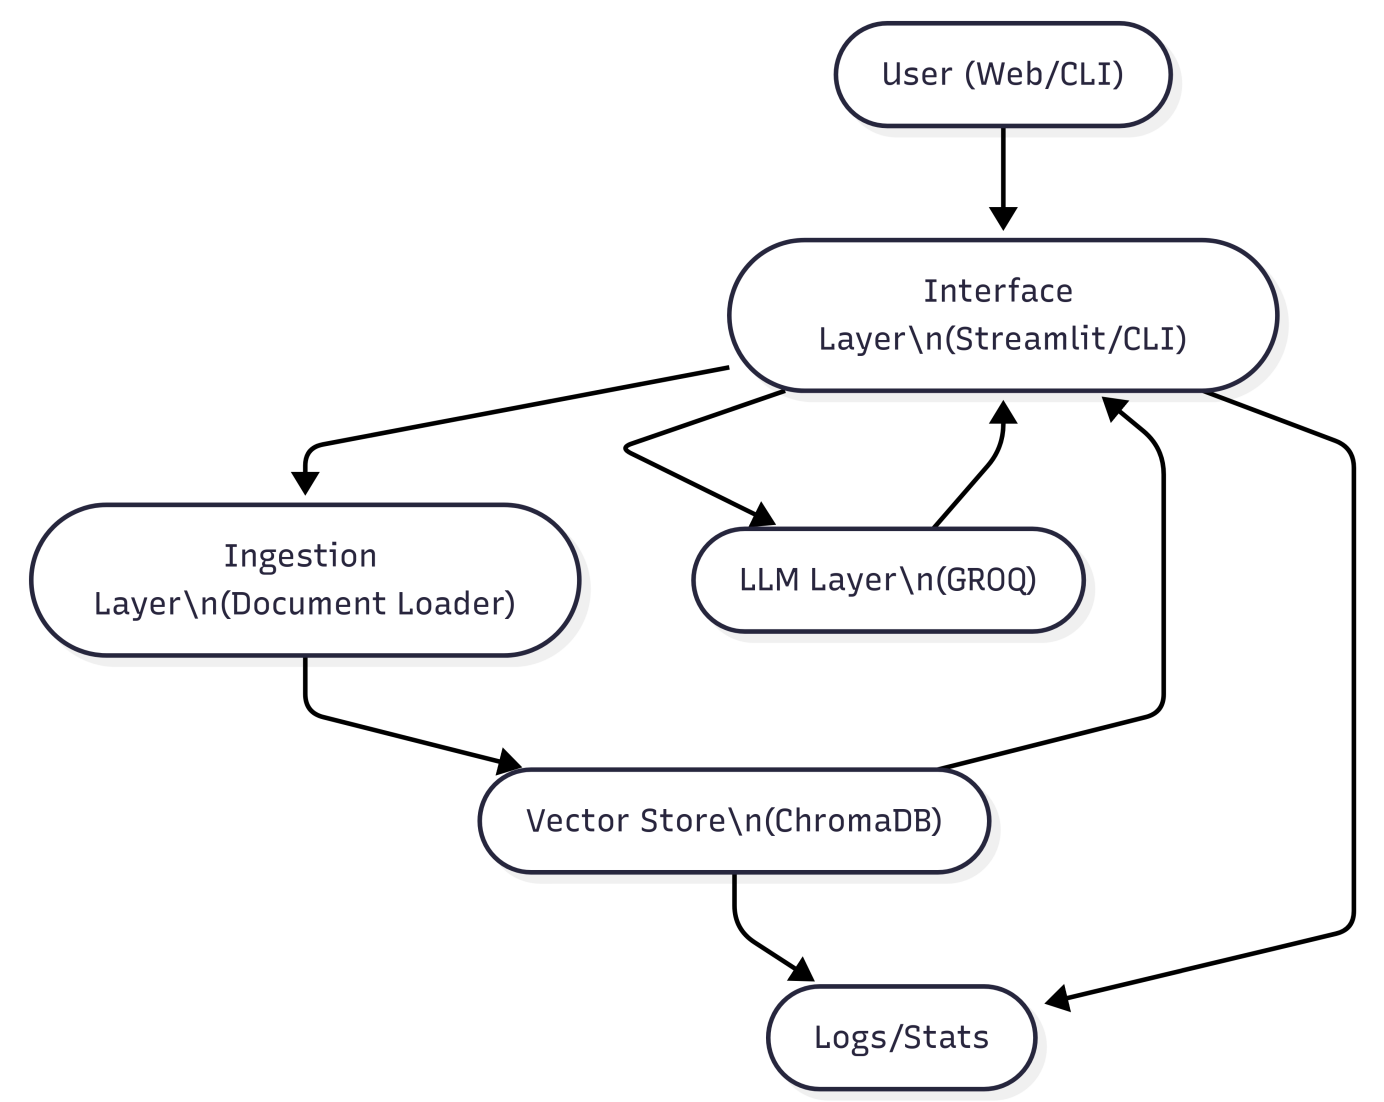


Figure 2: System Architecture showing the main components and their interactions

3. Key Modules and Their Roles

- **main.py:** CLI entry point; orchestrates all backend operations.
- **app.py:** Streamlit web interface; mirrors CLI functionality with a modern UI.
- **src/rag_pipeline.py:** Core pipeline; manages ingestion, retrieval, and LLM calls.
- **src/ingestion/document_loader.py:** Loads and parses PDF, DOCX, TXT, and JSON files.
- **src/utils/config_loader.py:** Loads YAML config with dot-notation access.
- **src/utils/init_manager.py:** Initializes logging, loads .env, and sets up environment.
- **src/utils/log_manager.py:** Handles timestamped log file creation and management.

4. Data Flow and Processing Pipeline

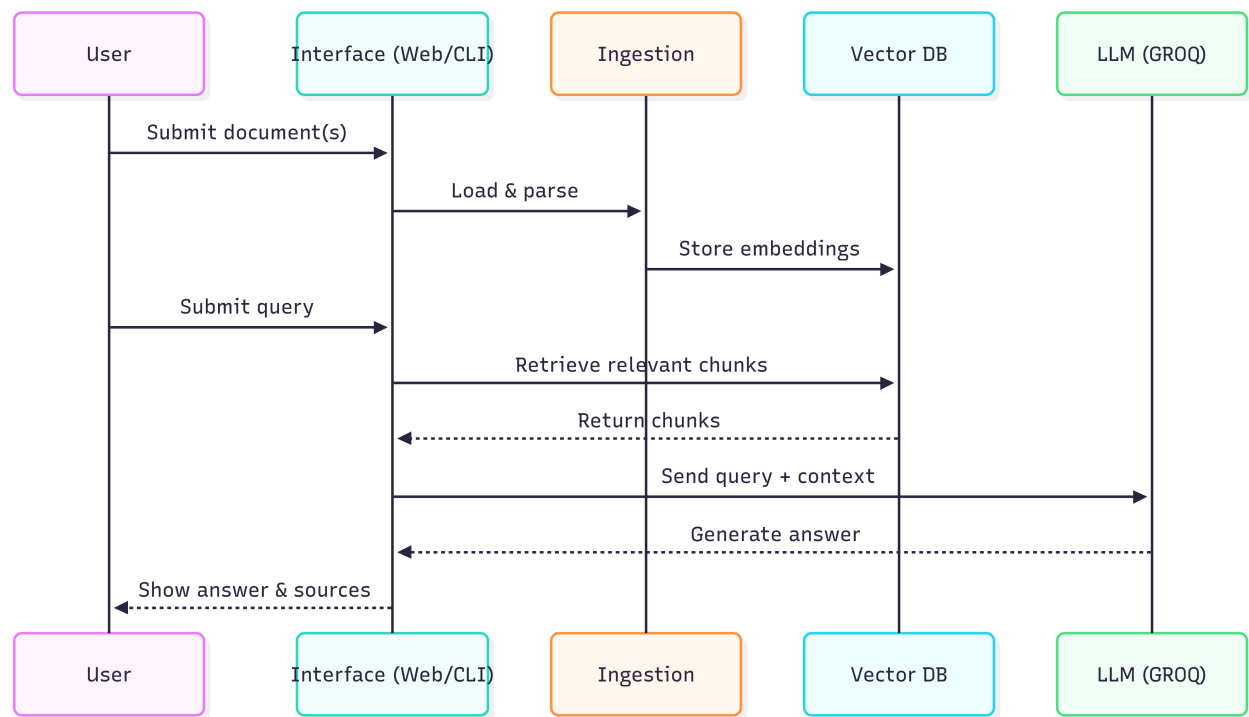


Figure 3: Data flow sequence showing the end-to-end process from document ingestion to query response

5. Security and Configuration

- **Secrets Management:** All API keys (e.g., GROQ_API_KEY) are loaded from a `.env` file (see `.env.example`).
- **Configuration:** System behavior is controlled via `config/config.yaml` (logging, LLM, vector DB, etc.).
- **Best Practices:** `.env` is git-ignored; `.env.example` is provided for onboarding.

6. Extensibility and Customization

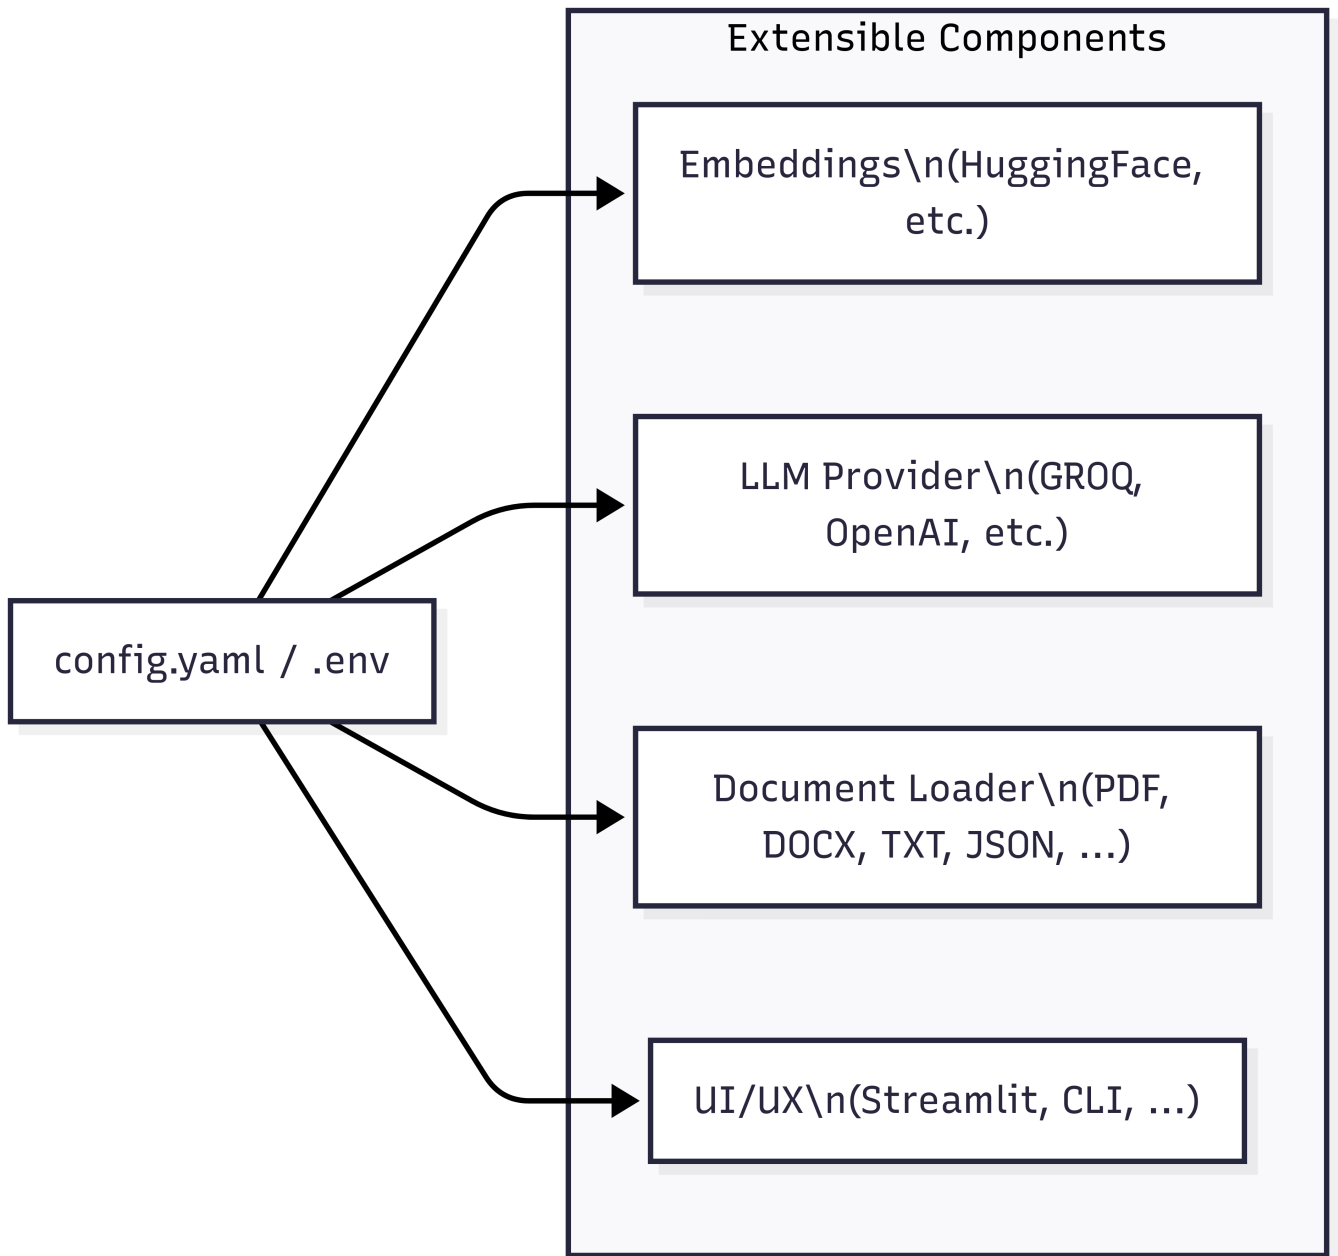


Figure 4: Extensibility diagram showing pluggable components and configuration-driven architecture

- **Pluggable Embeddings:** Swap HuggingFace models via config.
- **LLM Agnostic:** Easily switch LLM providers by updating config and .env.
- **Custom Ingestion:** Extend `document_loader.py` for new file types.
- **UI/UX:** Add new Streamlit pages or CLI commands as needed.

7. CLI and Web UI Design

- **CLI:** Supports all operations (init, ingest, query, stats, clear, logs, etc.) with rich help and examples.
- **Web UI:** Streamlit app with dashboard, chat, ingestion, stats, and log management.
- **Interactive Mode:** CLI supports conversational Q&A with `/stats`, `/help`, `/quit` commands.

8. Logging, Monitoring, and Testing

- **Logging:** Timestamped log files per session; configurable via YAML.
- **Monitoring:** Real-time stats in both CLI and web UI.

- **Testing:** CLI commands for component and end-to-end tests; pytest integration.

9. Supported Formats and Deployment

- **Documents:** PDF, DOCX, TXT, JSON (extensible).
- **Deployment:** Cross-platform (Windows, Linux, Mac); web UI via Streamlit; CLI via Python.
- **Hosting:** See [docs/Streamlit_Hosting.md](#) for deployment options.

10. References and Future Work

- **References:**
 - ChromaDB: <https://www.trychroma.com/>
 - LangChain: <https://www.langchain.com/>
 - GROQ: <https://groq.com/>
 - Streamlit: <https://streamlit.io/>
- **Future Work:**
 - Add support for more LLM providers (OpenAI, Azure, etc.)
 - Advanced analytics and visualization modules
 - Distributed/clustered vector store support
 - Enhanced security and RBAC for multi-user deployments
 - Automated document ingestion pipelines

For implementation details, see the codebase and README. For deployment, see the Streamlit hosting guide in [docs/](#).